

OpenTDX: Emulating TDX Machine

Jaewon Hur

SSLAB Georgia Tech

Supervisor: Taesoo Kim

Agenda

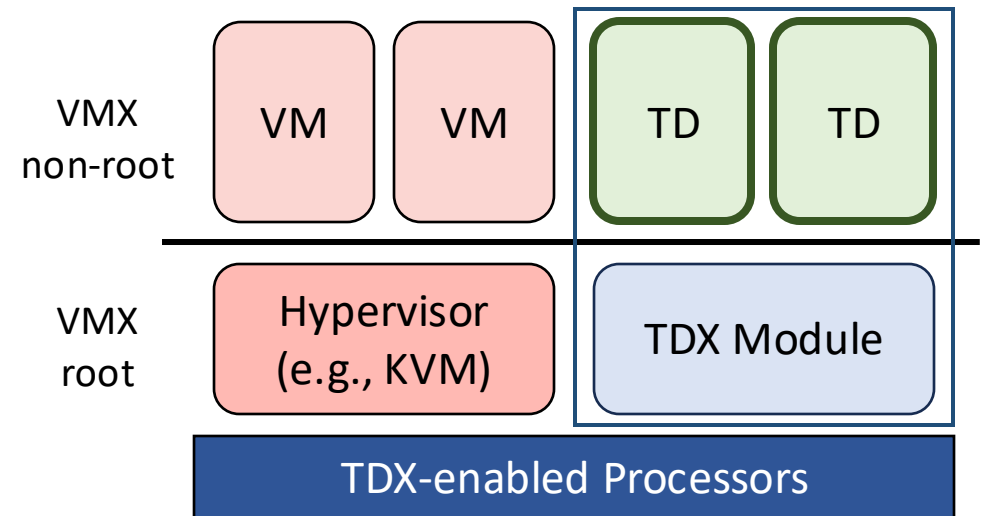
- **Introduction of Intel TDX**
 - **Motivation of OpenTDX**
- **Demo of Working Example**
- **Use Cases of OpenTDX**
- **Technical Details**
- **Conclusion**

What is TDX?

- **Intel TDX (Trust Domain Extensions)**
 - Set of HW & SW extensions for **VM based confidential computing**
 - Protect user's workloads on possibly compromised cloud environments
 - Open new data applications by ***guaranteeing the security of data***
 - E.g., medical dataset analysis on public cloud → regulation compliance
 - Convince end users that their private data is securely processed
 - E.g., secure LLM serving system

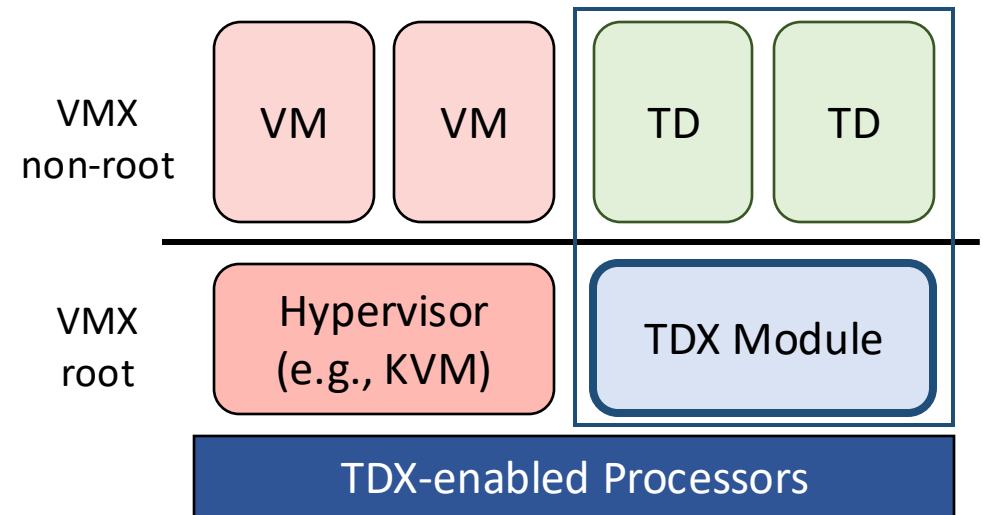
What is TDX?

- **Intel TDX (Trust Domain Extensions)**
 - Set of HW & SW extensions for **VM based confidential computing**
 - **TDs** (i.e., confidential VM) protected on cloud environment



What is TDX?

- **Intel TDX (Trust Domain Extensions)**
 - Set of HW & SW extensions for **VM based confidential computing**
 - **TDs** (i.e., confidential VM) protected on cloud environment
 - **TDX Module** (i.e., trusted thin hypervisor) protecting life-cycle of TDs



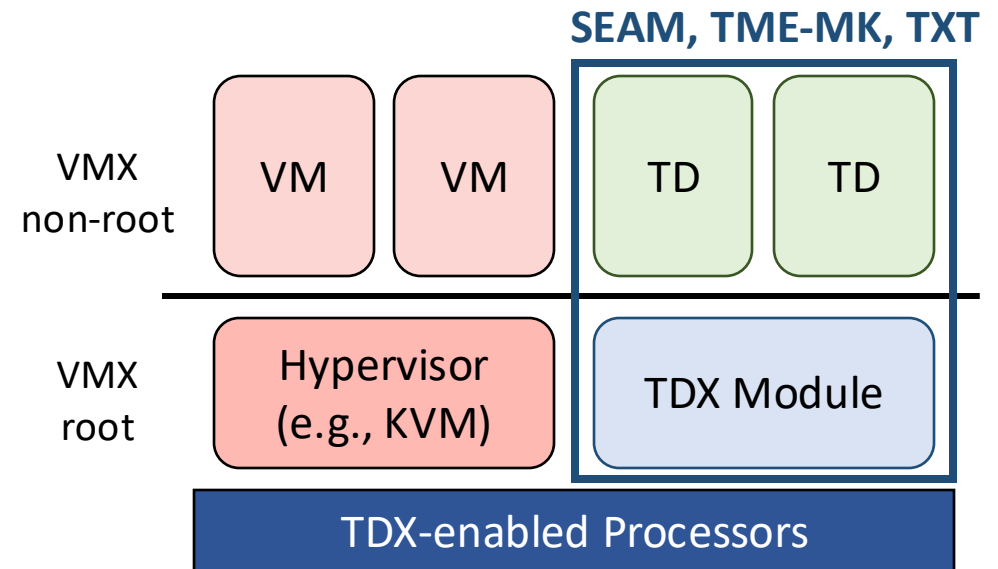
What is TDX?

- **Intel TDX (Trust Domain Extensions)**

- Set of HW & SW extensions for **VM based confidential computing**
- **TDs** (i.e., confidential VM) protected on cloud environment
- **TDX Module** (i.e., trusted thin hypervisor) protecting life-cycle of TDs

- Hardware extensions

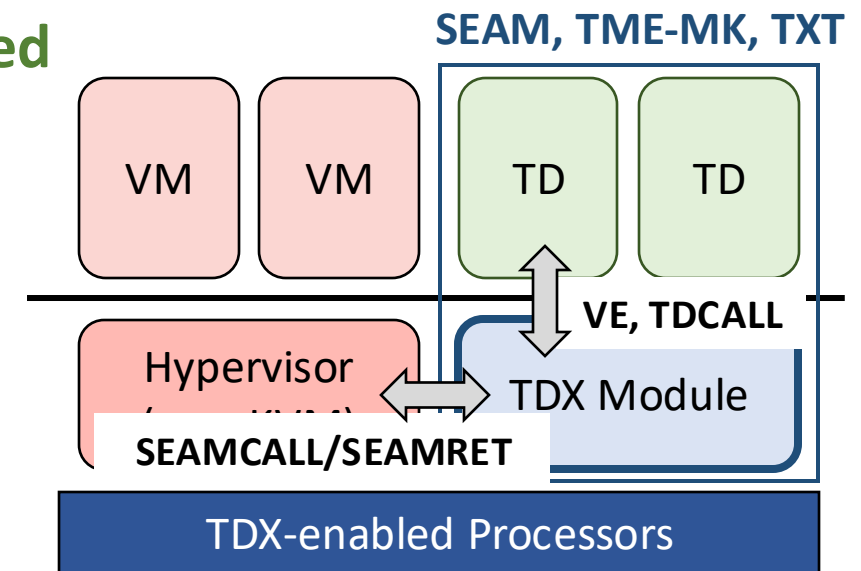
- **SEAM** (SEcure Arbitration Mode)
 - For access control in CPU
- **TME-MK** (Total Memory Encryption Multi-Key)
 - For encrypting data at memory
- **TXT** (Trusted Execution Technology)
 - For root-of-trust



What is TDX Module?

- **TDX Module**

- Trusted thin hypervisor for constructing, operating, destructing TDs
 - Emulate CPU instructions, manage memories, etc.
- Interacts with host hypervisor (i.e., **SEAMCALL/SEAMRET**), and TDs (i.e., **Virtualization Exception, TDCALL**)
- **Only the TDX Module signed by Intel can be loaded**
 - Root of Trust for TDs
- **Unable to customize TDX Module**
 - Limited contributions from research community
 - Difficulty of implementing and debugging TDX SW stacks

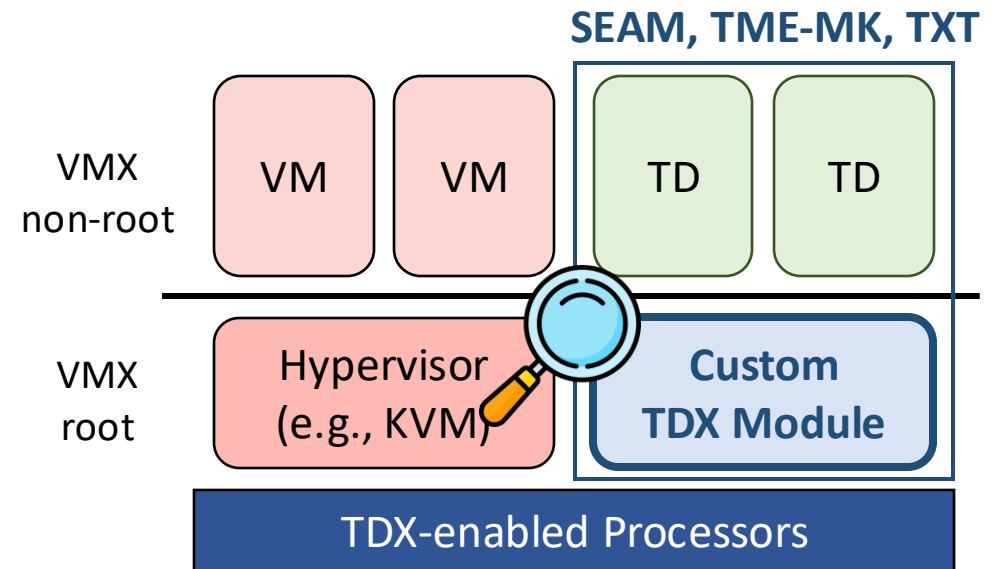


Bridging the Gap: OpenTDX

- **OpenTDX**

- We emulate a TDX machine (on a normal Intel processor) to run **custom TDX modules**
- We extend **QEMU-KVM** to emulate a TDX machine

- **Any one can develop and test their own TDX modules on OpenTDX**



Demo: Working Example

Project Description

- <https://github.com/sslab-gatech/open-tdx>

Use Cases of OpenTDX

- Re-implementing TDX Module in RUST
- Secure & efficient shared memory between TDs
 - TDX Module can assign the same shared TME-MK key for the TDs
- Fuzzing TDX Module (i.e., white-box fuzzing)
 - Instrument TDX Module for sanitization, coverage guidance, etc.
 - [Cornelius \[1\]](#) can be used also, but some functionalities would be limited (e.g., VE injection)
- Implementing & debugging software stack for TDX
 - New host hypervisor
 - New TD kernel

[1] <https://github.com/microsoft/Cornelius>

Use Cases of OpenTDX

- New Design of Secure LLM Serving System using TDX Module

- Motivation

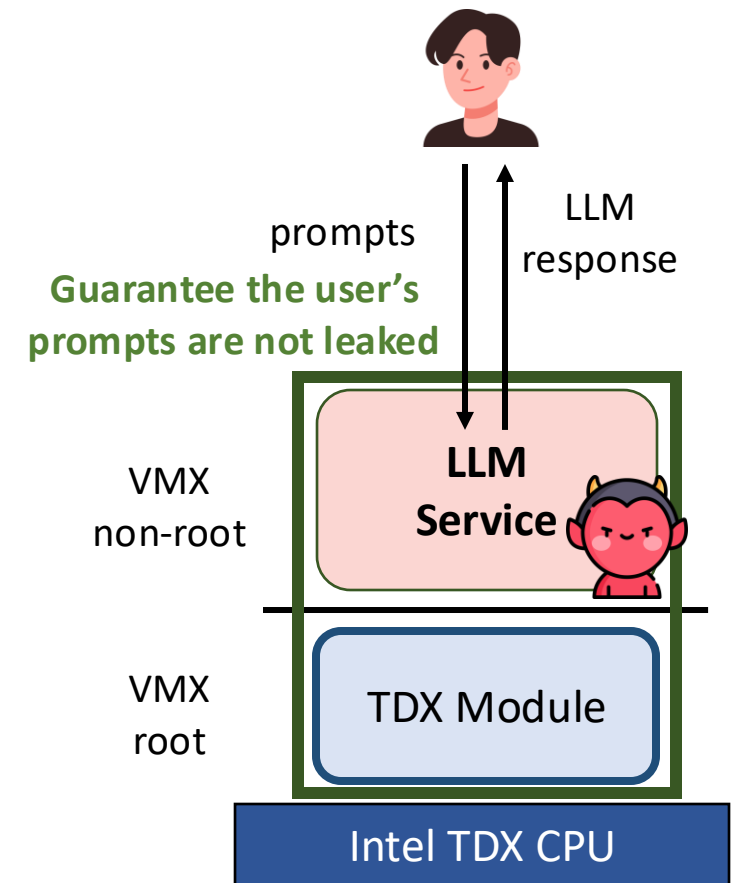
- LLM service may be compromised (or even malicious) due to its large code base

- Goal

- Protect the user's prompts even when the LLM service is compromised

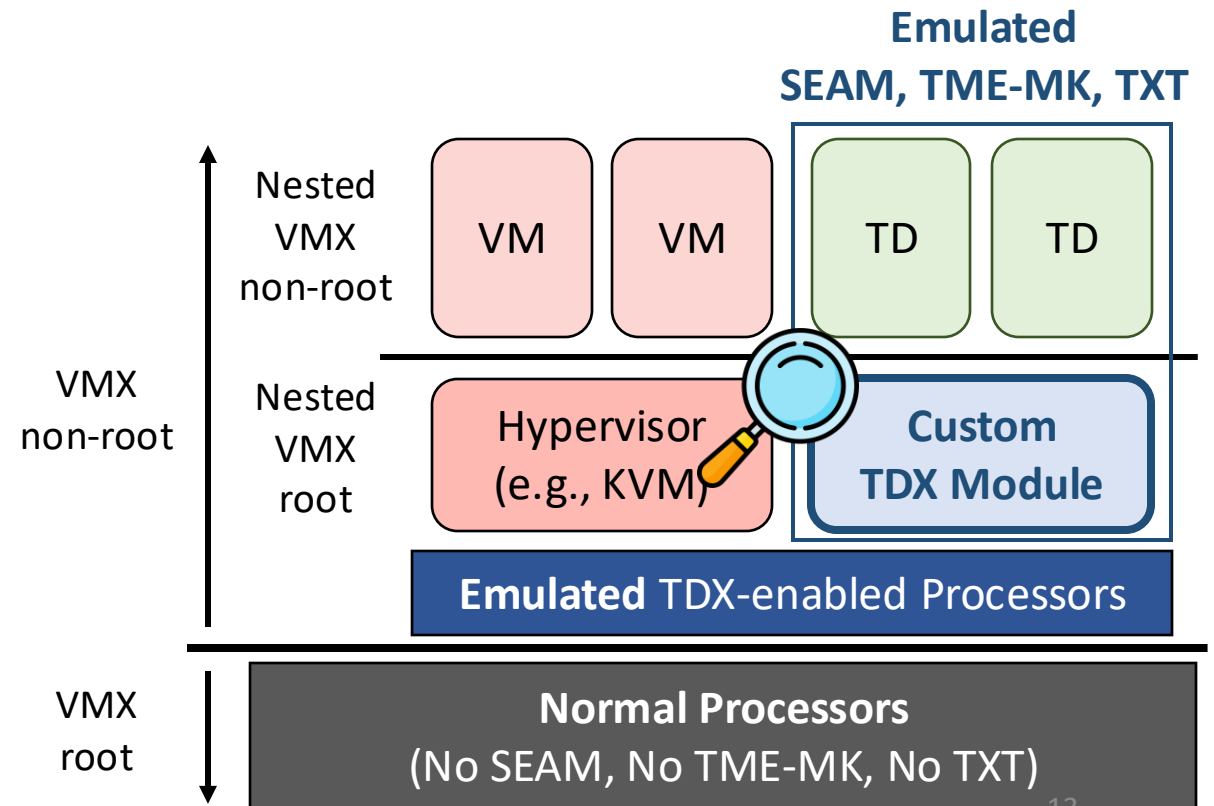
- Key Insight

- TDX module can provide a stronger protection on the prompts thanks to its higher privilege
 - TDX module can be open-sourced and verified by end users



Deep Dive into OpenTDX

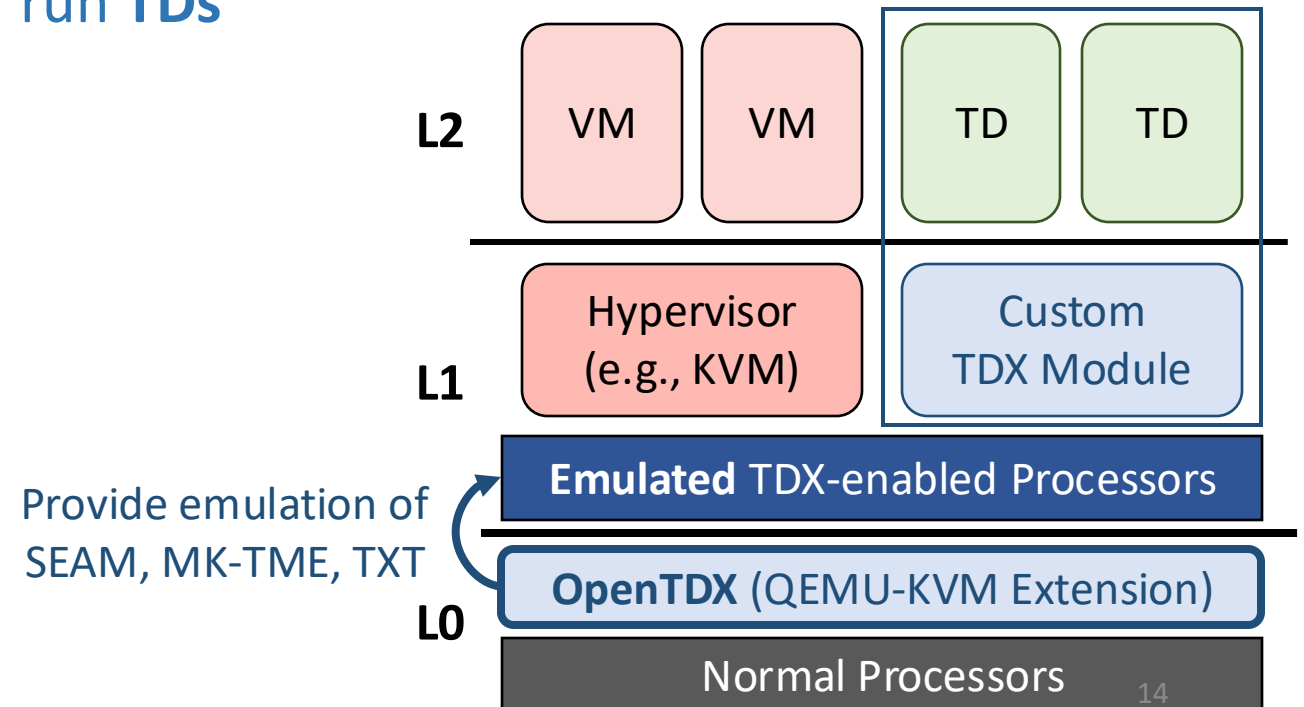
- Components of OpenTDX
 - Nested virtualization
 - Host TDs as nested VMs
 - ISA emulation
 - TXT
 - TME-MK
 - SEAM
 - Other TDX features



Nested Virtualization for Serving TDs

- Virtualization Levels

- **L0:** Normal Intel CPU (i.e., no TDX support) → run **OpenTDX**
- **L1:** Emulated TDX CPU → run **TDX Module**
- **L2:** VMs on emulated TDX CPU → run **TDs**



ISA Emulation: TXT

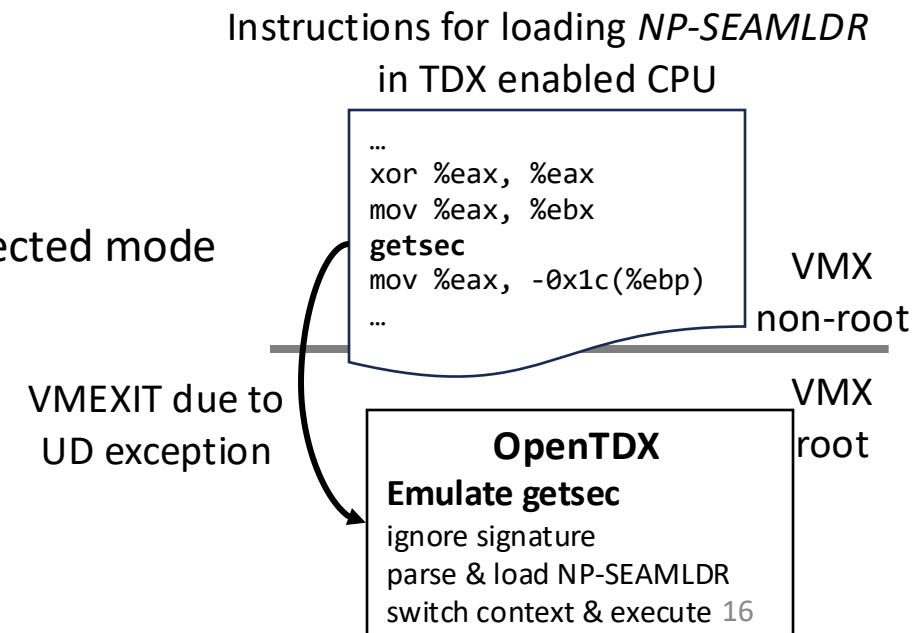
- Trusted Loading Procedure of TDX Module
 - ***NP-SEAMLDR*** → ***P-SEAMLDR*** → ***TDX Module***
 - Intel TXT guarantees only the ***NP-SEAMLDR*** signed by Intel is loaded
 - ***NP-SEAMLDR*** further guarantees only the correct ***TDX Module*** is loaded
- ISA
 - GETSEC[ENTERACCS]
 - Given a signed NP-SEAMLDR binary, check signature, parse and load, then execute in a protected mode
 - GETSEC[EXITAC]
 - Exit from the protected mode

Instructions for loading *NP-SEAMLDR*
in TDX enabled CPU

```
...  
xor %eax, %eax  
mov %eax, %ebx  
getsec  
mov %eax, -0x1c(%ebp)  
...
```

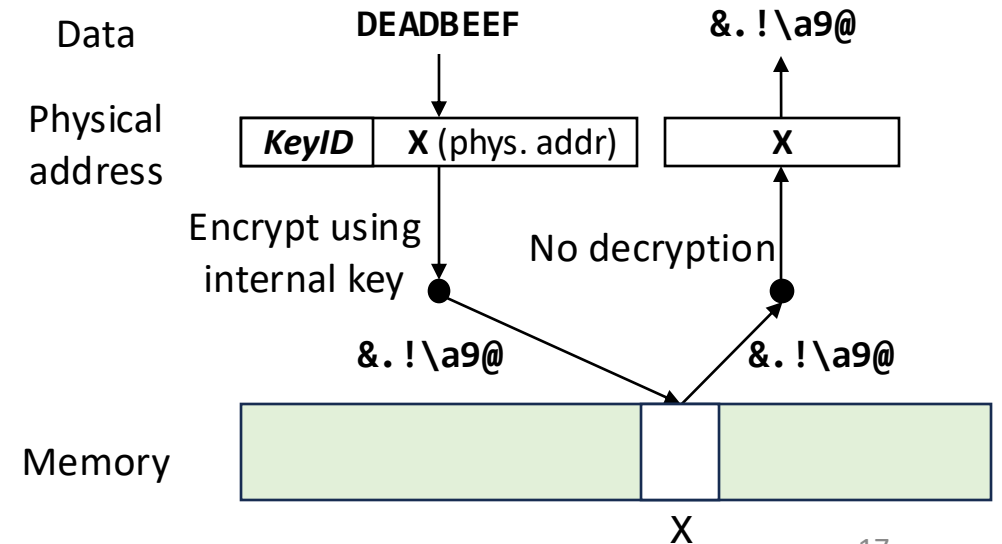
ISA Emulation: TXT

- Trusted Loading Procedure of TDX Module
 - NP-SEAMLDR* → *P-SEAMLDR* → *TDX Module*** **Removed signature verification**
 - Intel TXT guarantees only the ***NP-SEAMLDR*** ~~signed by Intel~~ is loaded
 - NP-SEAMLDR*** further guarantees only the correct ***TDX Module*** is loaded
- ISA **Emulate them by catching *UD* exception**
 - GETSEC[ENTERACCS]**
 - Given a signed NP-SEAMLDR binary, check signature, parse and load, then execute in a protected mode
 - GETSEC[EXITAC]**
 - Exit from the protected mode



ISA Emulation: TME-MK

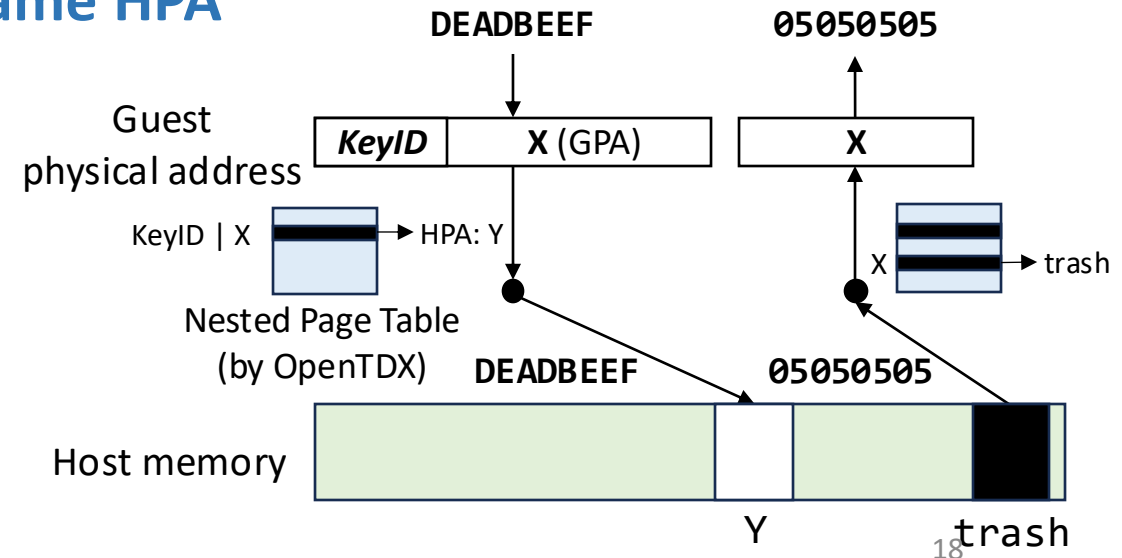
- Memory encryption of TDX module and TDs
 - CPU repurposes topmost bits of physical address as **KeyID**
 - Encrypts the outgoing data using **internal key bound to KeyID**
 - Read with different (or empty) **KeyID** returns trash value
 - TDX module protects memory by programming **KeyIDs**
- ISA
 - Memory encryption based on KeyID
 - TME-MK MSRs
 - Configure TME-MK keyID space
 - PCONFIG
 - Program KeyIDs



ISA Emulation: TME-MK

- Memory encryption of TDX module and TDs
 - CPU repurposes topmost bits of physical address as **KeyID** **Do not encrypt data**
 - ~~Encrypts the outgoing data using internal key bound to KeyID~~
 - Read with different (or empty) **KeyID** returns trash value
 - TDX module protects memory by programming **KeyIDs**
Manipulate nested page table to make (keyID | GPA) point to the same HPA
- ISA
(keyID | GPA) point to the same HPA
 - Memory encryption based on KeyID
 - TME-MK MSRs
 - Configure TME-MK keyID space
 - PCONFIG
 - Program KeyIDs

Emulate MSRs & instruction



ISA Emulation: SEAM

- Privilege Separation of TDX Module vs. Normal Hypervisors
 - CPU enters privileged mode (i.e., SEAM) upon executing **SEAMCALL**
 - **SEAMCALL** loads TDX module's context from **SEAM VMCS**
 - Only SEAM mode SW can access TME-MK encrypted data
 - TDX Module in SEAM mode serves TDs as like normal hypervisors
 - Using VT-X features (e.g., VMCS, VMENTER, etc.)
- ISA
 - Context switch using SEAM VMCS
 - Memory access control based on SEAM mode
 - Access permission to TME-MK encrypted memory
 - SEAM MSRs
 - For protecting SEAM memory regions

Instructions for entering TDX Module

Normal mode

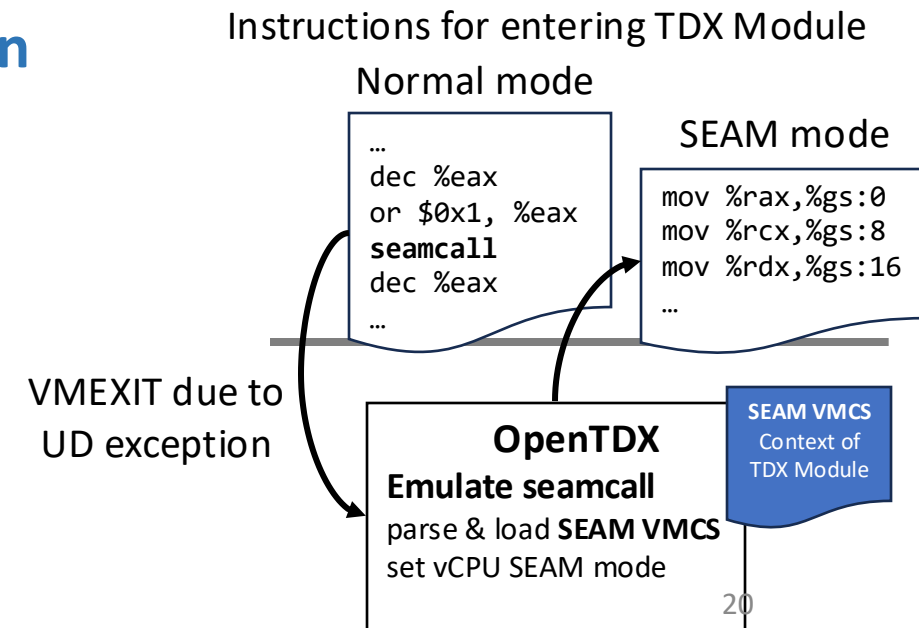
```
...  
dec %eax  
or $0x1, %eax  
seamcall  
dec %eax  
...
```

SEAM mode

```
mov %rax,%gs:0  
mov %rcx,%gs:8  
mov %rdx,%gs:16  
...
```

ISA Emulation: SEAM

- Privilege Separation of TDX Module vs. Normal Hypervisors
 - CPU enters privileged mode (i.e., SEAM) upon **executing SEAMCALL**
 - **SEAMCALL** loads TDX module's context from **SEAM VMCS** **Emulate SEAMCALL/SEAMRET**
 - Only SEAM mode SW can access TME-MK encrypted data
 - TDX Module in SEAM mode serves TDs as like normal hypervisors
 - Using VT-X features (e.g., VMCS, VMENTER, etc.)
Use nested virtualization
- ISA
 - Context switch using SEAM VMCS
Emulate VMCS structure
 - Memory access control based on SEAM mode
 - ~~Access permission to TME-MK encrypted memory~~
Currently not supported
 - SEAM MSRs
 - For protecting SEAM memory regions
Emulate MSRs



ISA Emulation: Other TDX Features

- GPA separation
 - Topmost bit = 0 in GPA → Normal EPT (encrypted using TD's KeyID)
 - Topmost bit = 1 in GPA → Shared EPT (non-encrypted)
 - Shared EPTP field added to VMCS

ISA Emulation: Other TDX Features

- GPA separation
 - On nested EPT violation, walk normal/shared EPT depending on the topmost bit
 - Topmost bit = 0 in GPA → Normal EPT (encrypted using TD's KeyID)
 - Topmost bit = 1 in GPA → Shared EPT (non-encrypted)
 - Shared EPTP field added to VMCS
 - Add field to nested VMCS

ISA Emulation: Other TDX Features

- GPA separation
 - On nested EPT violation, walk normal/shared EPT depending on the topmost bit
 - Topmost bit = 0 in GPA → Normal EPT (encrypted using TD's KeyID)
 - Topmost bit = 1 in GPA → Shared EPT (non-encrypted)
 - Shared EPTP field added to VMCS
 - Add field to nested VMCS
- Virtualization Exception
 - EPT violation & specific EPT entry setting
 - Inject virtualization exception in TD without VMEXIT

ISA Emulation: Other TDX Features

- GPA separation
 - On nested EPT violation, walk normal/shared EPT depending on the topmost bit
 - Topmost bit = 0 in GPA → Normal EPT (encrypted using TD's KeyID)
 - Topmost bit = 1 in GPA → Shared EPT (non-encrypted)
 - Shared EPTP field added to VMCS
 - Add field to nested VMCS
- Virtualization Exception
 - On nested EPT violation, check condition & inject
 - EPT violation & specific EPT entry setting
→ Inject virtualization exception in TD without VMEXIT

Conclusion

- OpenTDX emulates TDX machines on normal Intel CPUs
- Wish you enjoy playing with OpenTDX!
- <https://github.com/sslabs-gatech/open-tdx>

Thank you

For any question, please contact jwhur19@gmail.com